

State, Immutability, and Persistent Data Structures

Brian Maddy
@bmaddy

Functional Relational Programming

(not reactive programming - see Colin Lee's talk for that)

Out of the Tar Pit

Ben Moseley
ben@moseley.name

Peter Marks
public@indigomail.net

February 6, 2006

Abstract

Complexity is the single major difficulty in the successful development of large-scale software systems. Following Brooks we distinguish *accidental* from *essential* difficulty, but disagree with his premise that most complexity remaining in contemporary systems is essential. We identify common causes of complexity and discuss general approaches which can be taken to eliminate them where they are accidental in nature. To make things more concrete we then give an outline for a potential complexity-minimizing approach based on *functional pro-*

The Plan

- Amount and impact of mutable state
- Identity vs. Value
- Benefits of less mutable state
- What can we do now?

Sound familiar?

“Try it again”

“Reboot”

“Restart the program”

“It must have been a fluke”

“Try refreshing the page”

“Restart the server”

“Reinstall the program”

“Well, it seems to work now”

“Reload the file”

“Reinstall the operating system and then the program”

How many possible states?

1 bit: 2 possible states

How many possible states?

1 bit: 2 possible states

of atoms on Earth: 2^{167} atoms

How many possible states?

1 bit: 2 possible states

of atoms on Earth: 2^{167} atoms

6 32-bit long ints: 2^{192} possible states

How many possible states?

1 bit: 2 possible states

of atoms on Earth: 2^{167} atoms

6 32-bit long ints: 2^{192} possible states

of atoms in the Universe: 2^{266} atoms

How many possible states?

1 bit: 2 possible states

of atoms on Earth: 2^{167} atoms

6 32-bit long ints: 2^{192} possible states

of atoms in the Universe: 2^{266} atoms

5 JavaScript numbers: 2^{320} possible states

How many possible states?

1 bit: 2 possible states

of atoms on Earth: 2^{167} atoms

6 32-bit long ints: 2^{192} possible states

of atoms in the Universe: 2^{266} atoms

5 JavaScript numbers: 2^{320} possible states

8GB of RAM: $2^{64000000000}$ possible states

Program Correctness

$$P = p^n$$

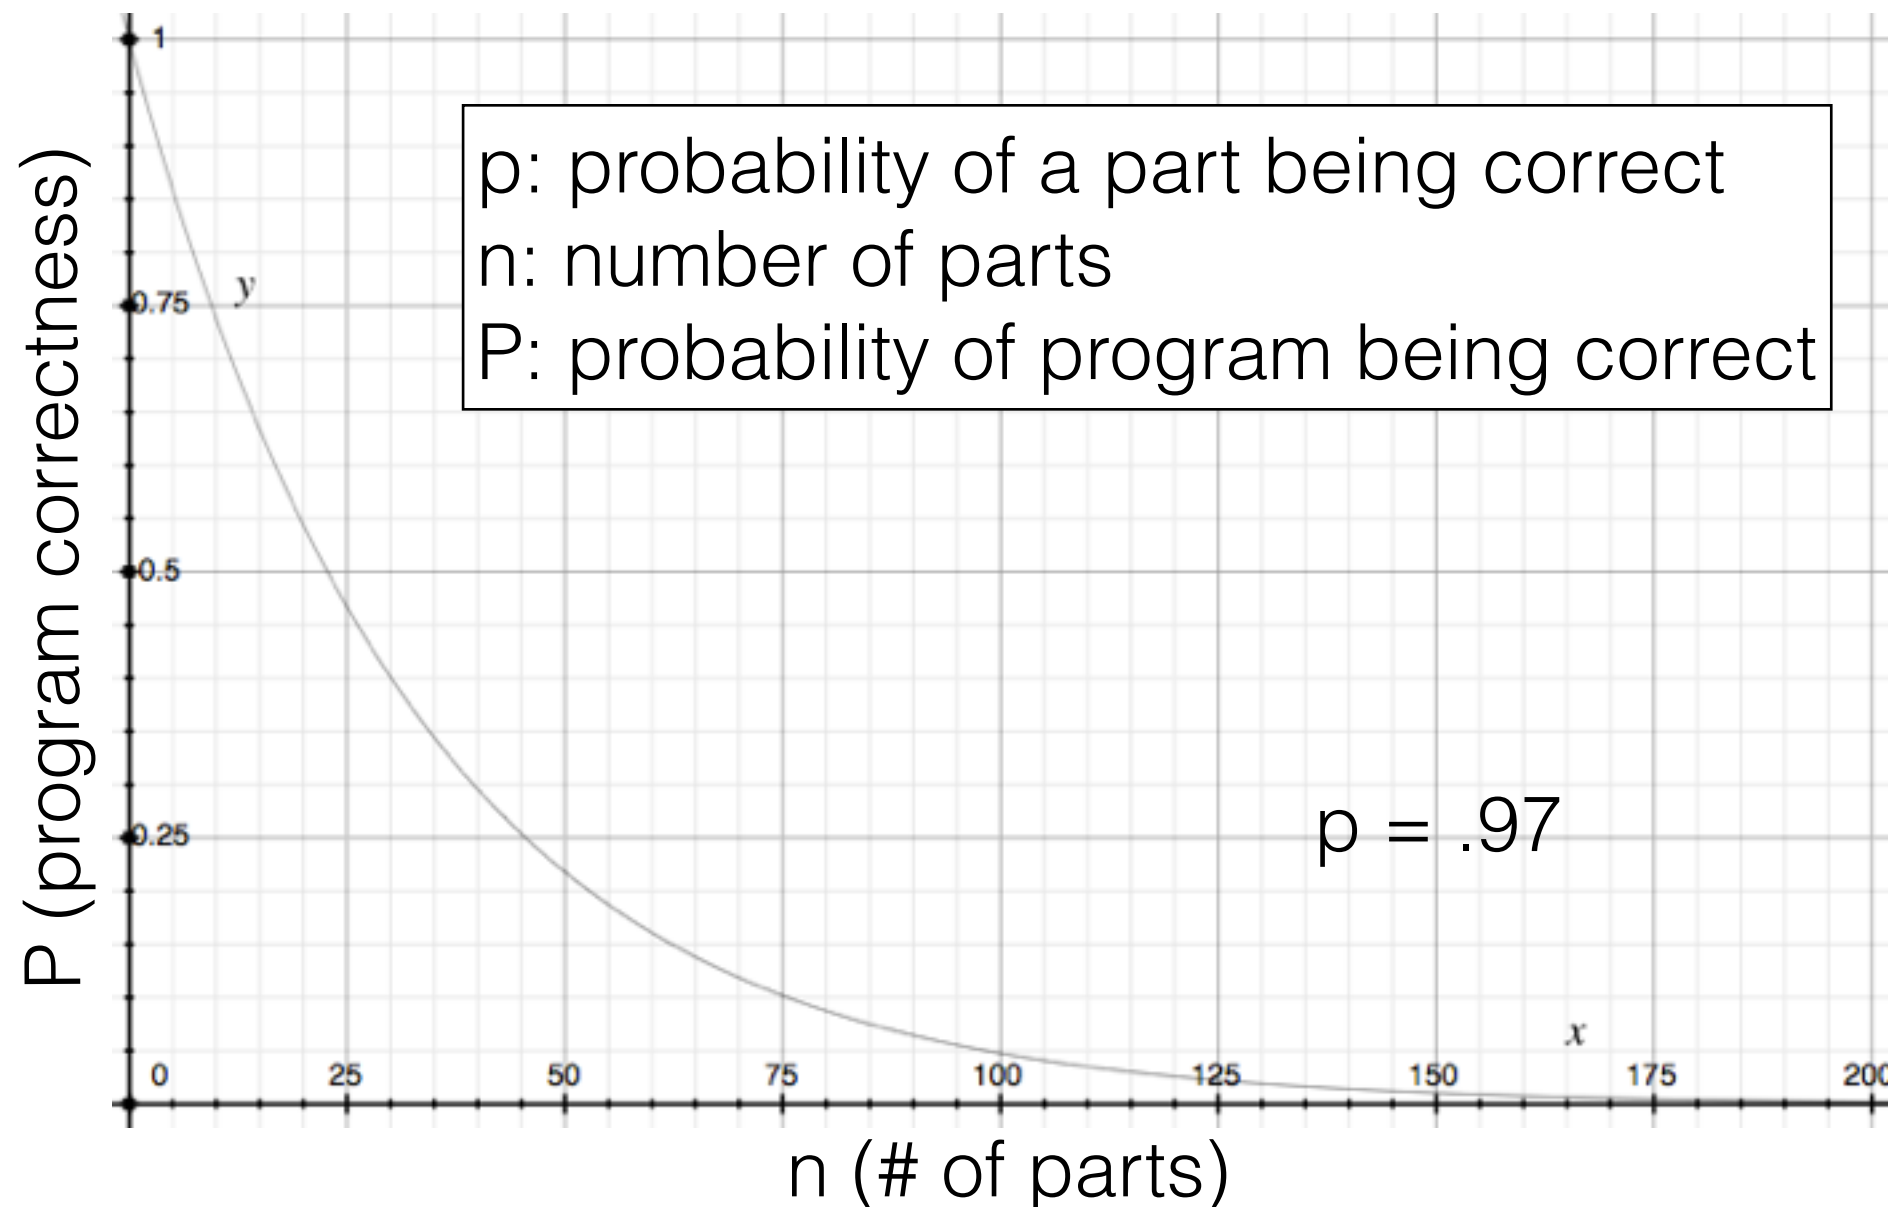
p: probability of a part being correct

n: number of parts

P: probability of program being correct

Program Correctness

$$P = p^n$$



Pure Functions

- always returns the same result for a given input
(uses no mutable state)
- no side effects

Mutable State Pollutes

```
(defn stateless-func []  
  ...do stuff...  
  ...do stuff...  
  ...do stuff...)
```

Mutable State Pollutes


Not anymore!





```
(defn stateless-func []  
  ...do stuff...  
  (stateful-func)  
  ...do stuff...)
```

Concurrency

Variable double(x) triple(x)




Concurrency




Concurrency


Variable double(x) triple(x)



Variable double(x) triple(x)




Variable double(x) triple(x)




Concurrency


Variable double(x) triple(x)




Variable double(x) triple(x)



Variable double(x) triple(x)




Variable double(x) triple(x)




Concurrency


Variable double(x) triple(x)




Variable double(x) triple(x)




Variable double(x) triple(x)



Variable double(x) triple(x)




Variable double(x) triple(x)




Concurrency


Variable double(x) triple(x)




Variable double(x) triple(x)




Variable double(x) triple(x)




Variable double(x) triple(x)



Variable double(x) triple(x)




Variable double(x) triple(x)




Concurrency


Variable double(x) triple(x)




Variable double(x) triple(x)




Variable double(x) triple(x)




Variable double(x) triple(x)



Variable double(x) triple(x)



Variable double(x) triple(x)



“testing is hopelessly inadequate”

– Edgar Dijkstra, *Notes on Structured Programming*

Get rid of all state?

Mutable Counter

```
procedure int getNextCounter()  
    // counter is initialized elsewhere  
    count = count + 1  
    return count
```

// Usage:

```
getNextCounter() // => 1  
getNextCounter() // => 2  
getNextCounter() // => 3
```

Immutable Counter

```
function (int,int) getNextCounter(int oldCount)
  int result = oldCount + 1
  int newCount = oldCount + 1
  return (newCount, result)
```

// Usage:

```
countA, resultA = getNextCounter(0)
countB, resultB = getNextCounter(countA)
countC, resultC = getNextCounter(countB)
resultC // => 3
```


Sometime mutable
state is nice...

How can we minimize
mutable state?

Essential State vs. Derived State

Identity vs. Value

Identity is a succession of values



Versioned Backups

(code)

“This, milord, is my family's axe. We have owned it for almost nine hundred years, see. Of course, sometimes it needed a new blade. And sometimes it has required a new handle, new designs on the metalwork, a little refreshing of the ornamentation ... but is this not the nine hundred-year-old axe of my family? And because it has changed gently over time, it is still a pretty good axe, y'know. Pretty good.”

– Terry Pratchett, *The Fifth Elephant*

“No man can cross the same river twice,
because neither the man nor the river are the
same.”

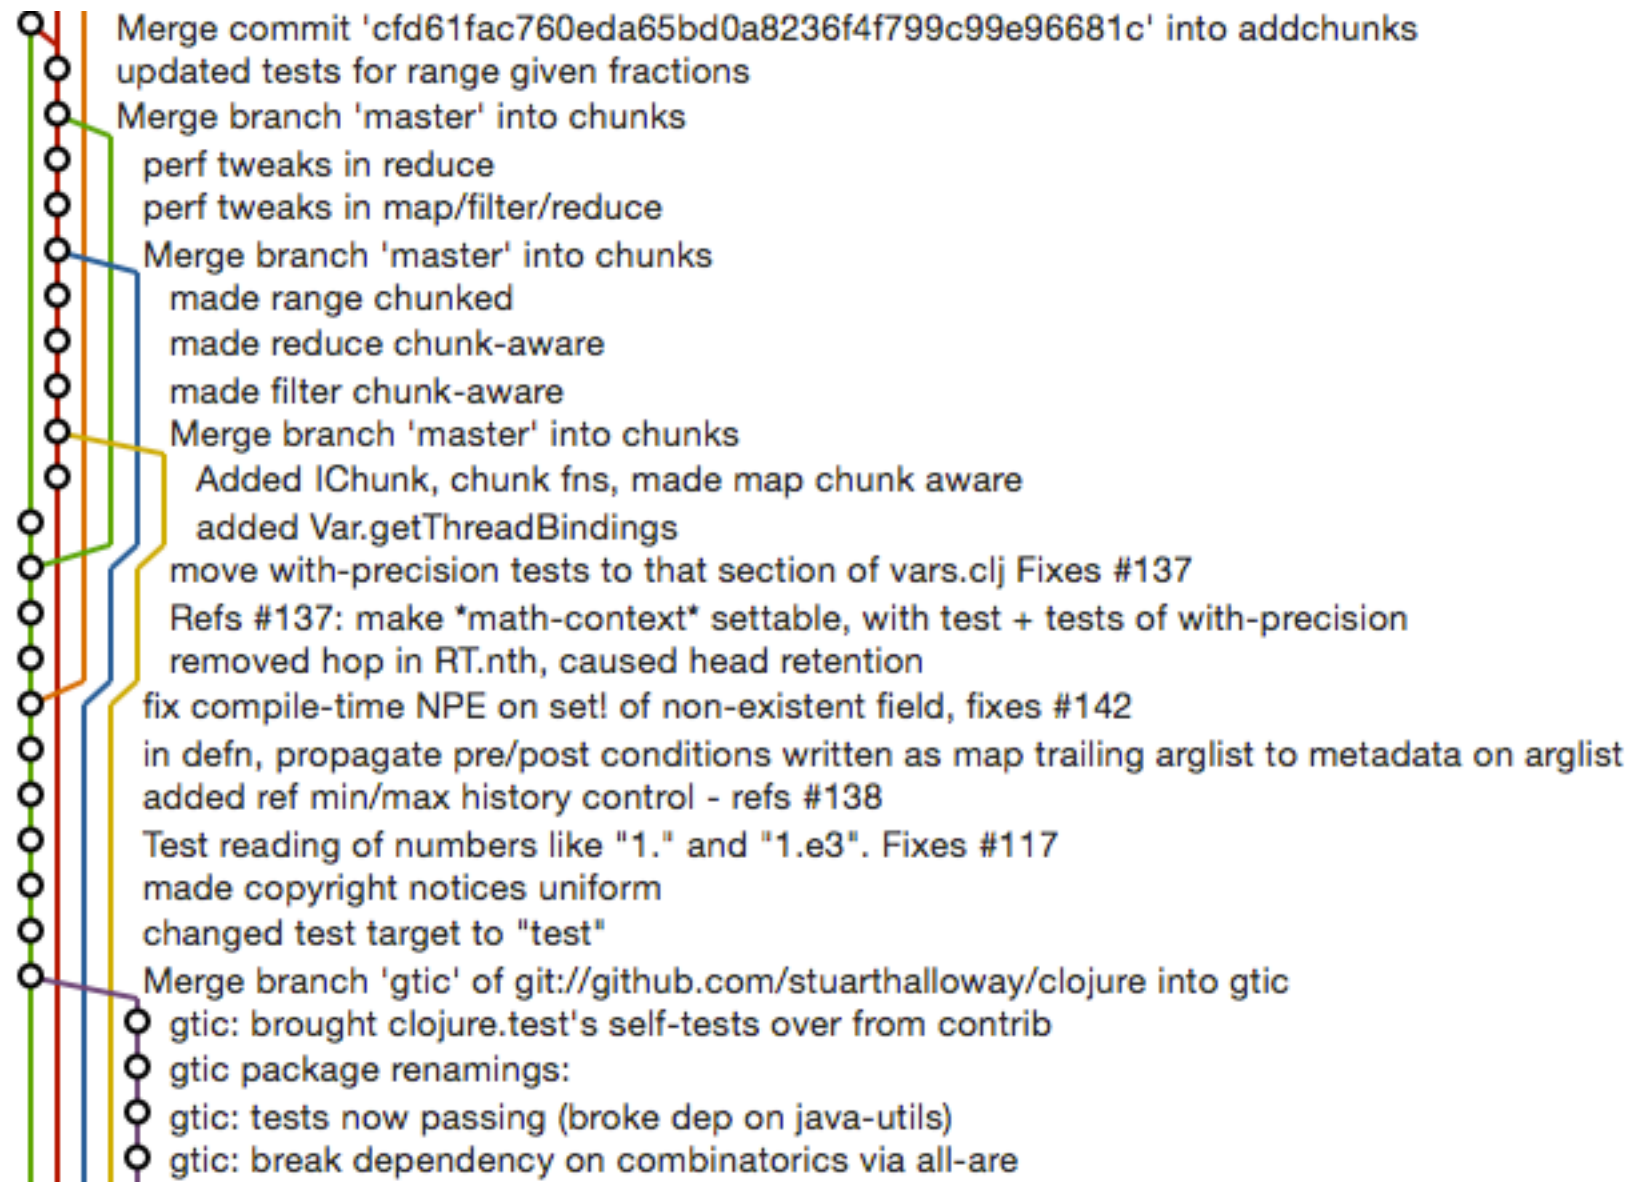
–Heraclitus

(code)

“[Think] of an experience from your childhood. Something you remember clearly, something you can see, feel, maybe even smell, as if you were really there. After all you really were there at the time, weren't you? How else could you remember it? But here is the bombshell: you weren't there. Not a single atom that is in your body today was there when that event took place. Every bit of you has been replaced many times over (which is why you eat, of course). You are not even the same shape as you were then. The point is that you are like a cloud: something that persists over long periods, while simultaneously being in flux. Matter flows from place to place and momentarily comes together to be you. Whatever you are, therefore, you are not the stuff of which you are made.”

—Steve Grand, *Creation: Life and How to Make It*


Git




Let's make identities that
point to successive
values!

But deep copying is slow
and takes up a lot of
space...

Persistent Data Structures



Persistent Data Structures



Benefits

Benefits

- Concurrency: can share state

Benefits

- Concurrency: can share state
- Instant history

Benefits

- Concurrency: can share state
- Instant history
- No defensive copying

Benefits

- Concurrency: can share state
- Instant history
- No defensive copying
- `==` vs. `.equals()`

What can be done now?

1. Write pure functions when possible (don't modify state, only return new versions)
2. Use `map`, `reduce`, `filter`, etc. instead of `for` loops
3. Use persistent data structures when possible
4. Differentiate between inherent state and derived state (make derived state immutable)
5. Highly performance sensitive code? Resort to mutable state

(demo)

Thanks!

@bmaddy

Resources

<http://clojure.mn> - local Clojure user group

<http://4clojure.org> - Koan style Clojure puzzles

<http://vidku.com> - We're hiring - Front end, Back end,
iOS, and Android - postings on tech.mn